

Convolutional Neural Network (CNN)

Prof. Ph.D. Woo Youn Kim Chemistry, KAIST



Goals

5주	주제	Multilayer perceptron 2					
	목표	Knowing various issues on MLP and techniques to resolve them					
	내용	Overfitting, regularization, dropout, batch normalization, cross validation					
	주제	Convolutional Neural Network (CNN) & SMILES					
6주	목표	Understanding CNN and molecular representation with SMILES					
~ T	내용	Convolution, receptive field, stride, pooling Supervised learning of Log P and TPSA					
	주제	Molecular graphs & Graph Neural Network (GNN)					
7주	목표	Understanding GCN and molecular representation with graphs					
	내용	Molecular graph representation, graph convolutional network Supervised learning of logP and TPSA					



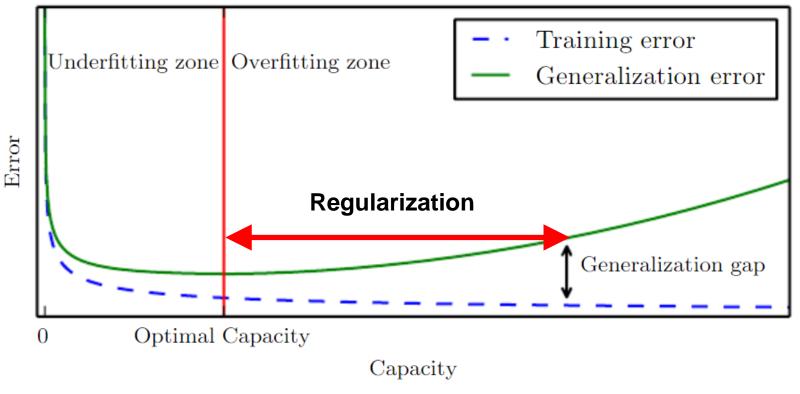
Contents

- Molecular representation with SMILES
- Convolutional neural network (CNN)
- Receptive field & Convolution
- Padding
- Skip-connection
- CNN in chemistry
- Practice session: supervised learning of log P and TPSA:



Review

- The main challenge is to find an optimal capacity of model for a given task.
- Regularization is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error.





CNN, RNN, GCN

Source

Deep Learning

An MIT Press book

lan Goodfellow and Yoshua Bengio and Aaron Courville

Exercises Lectures External Links

The Deep Learning textbook is a resource intended to help students and practitioners enter the field of machine learning in general and deep learning in particular. The online version of the book is now complete and will remain available online for free.

The deep learning textbook can now be ordered on Amazon.

https://www.deeplearningbook.org/

KAIST EE Jinwoo Shin http://alinlab.kaist.ac.kr/ee807 2018.html

http://cs231n.github.io/convolutional-networks/



Reference paper

Bioinformatics, 34, 2018, i821–i829 doi: 10.1093/bioinformatics/bty593 ECCB 2018



DeepDTA: deep drug-target binding affinity prediction

Hakime Öztürk¹, Arzucan Özgür^{1,*} and Elif Ozkirimli^{2,*}

¹Department of Computer Engineering and ²Department of Chemical Engineering, Bogazici University, Istanbul 34342, Turkey

*To whom correspondence should be addressed.



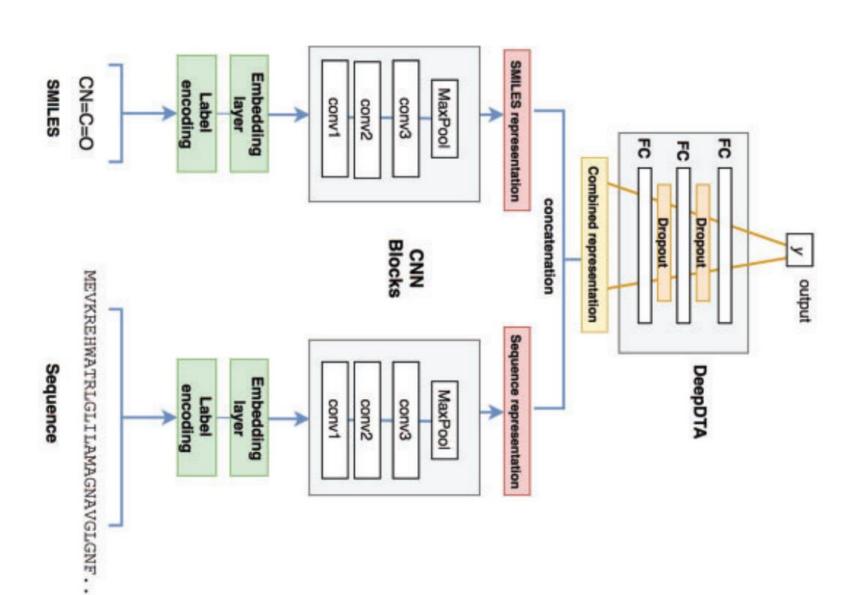
Reference paper

Abstract

Motivation: The identification of novel drug-target (DT) interactions is a substantial part of the drug discovery process. Most of the computational methods that have been proposed to predict DT interactions have focused on binary classification, where the goal is to determine whether a DT pair interacts or not. However, protein-ligand interactions assume a continuum of binding strength values, also called binding affinity and predicting this value still remains a challenge. The increase in the affinity data available in DT knowledge-bases allows the use of advanced learning techniques such as deep learning architectures in the prediction of binding affinities. In this study, we propose a deep-learning based model that uses only sequence information of both targets and drugs to predict DT interaction binding affinities. The few studies that focus on DT binding affinity prediction use either 3D structures of protein-ligand complexes or 2D features of compounds. One novel approach used in this work is the modeling of protein sequences and compound 1D representations with convolutional neural networks (CNNs).



Reference paper

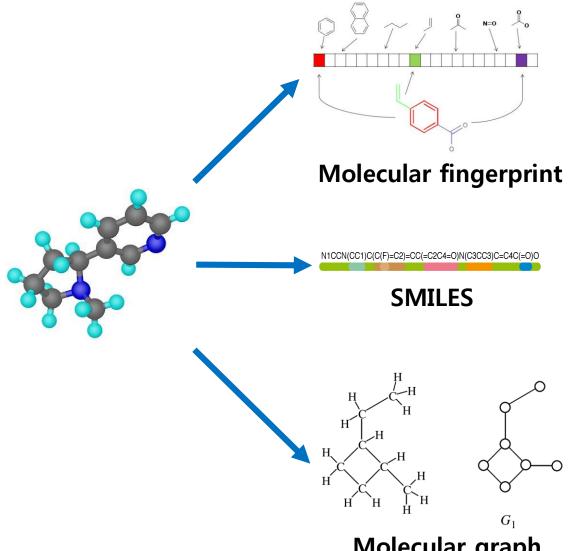




Molecular representation



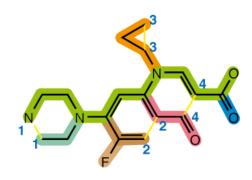
Molecular representation in deep learning





SMILES

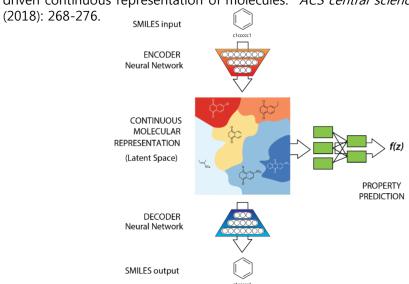
SMILES: 1D string



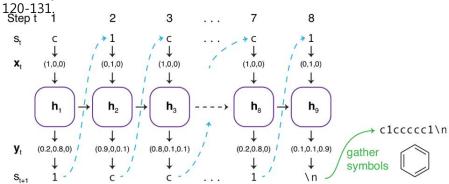
N1CCN(CC1)C(C(F)=C2)=CC(=C2C4=O)N(C3CC3)C=C4C(=O)O

- ✓ Simply match a molecule to string
- ✓ Easy to construct database
- ✓ Can use natural language AI models

Gómez-Bombarelli, Rafael, et al. "Automatic chemical design using a datadriven continuous representation of molecules." *ACS central science* 4.2



Segler, Marwin HS, et al. "Generating focused molecule libraries for drug discovery with recurrent neural networks." *ACS central science* 4.1 (2017):



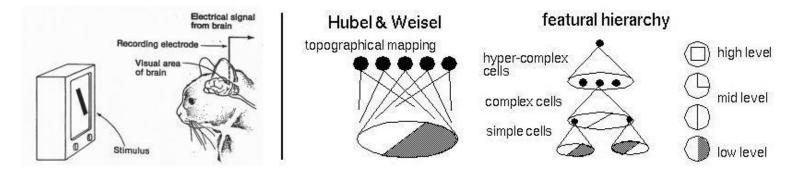


Introduction to CNN



History

Work by <u>Hubel</u> and <u>Wiesel</u> in the 1950s and 1960s showed that cat and monkey visual <u>cortexes</u> contain neurons that individually respond to small regions of the <u>visual field</u>.



The **receptive field** of an individual <u>sensory neuron</u> is the particular region of the sensory space (e.g., the body surface, or the visual field) in which a <u>stimulus</u> will modify the firing of that neuron.

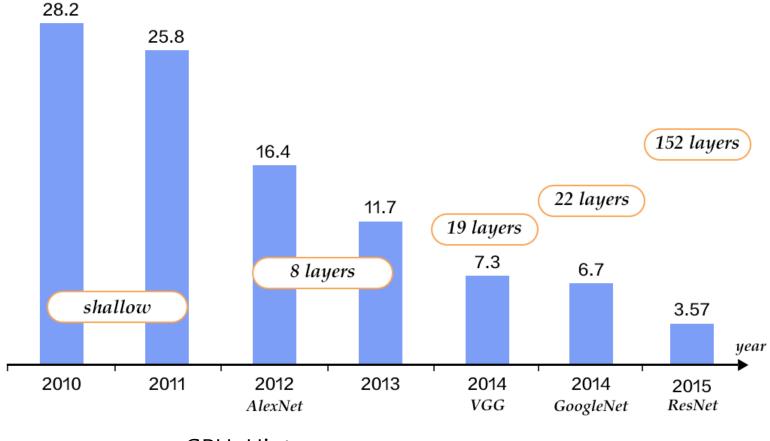
<u>LeCun</u> et al. in 1989, used back-propagation to learn the convolution kernel coefficients directly from images of hand-written numbers.

LeNet-5, a pioneering 7-level convolutional network by <u>LeCun</u> et al. in 1998



ILSVRC Winners

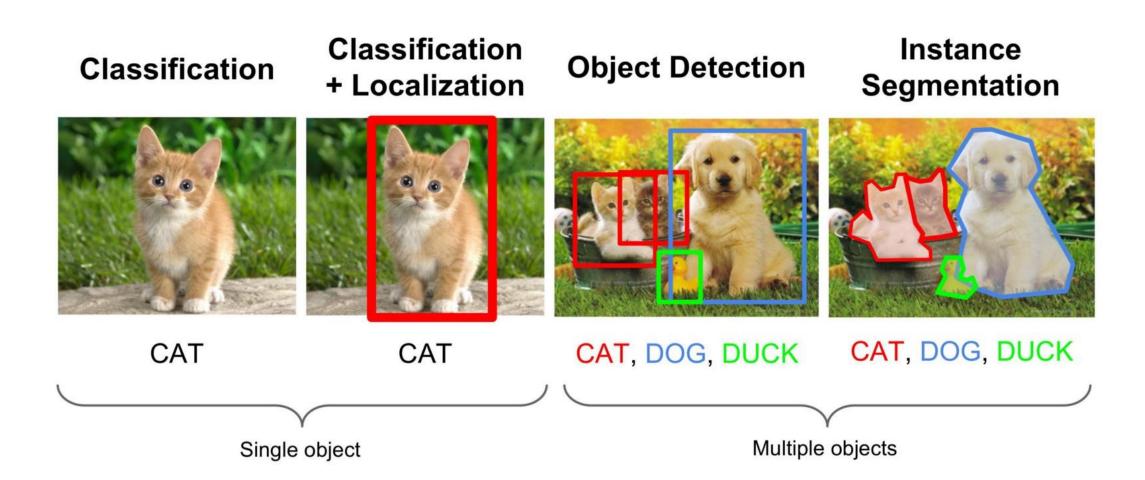
ImageNet Large Scale Visual Recognition Challenge





GPU, Hinton group

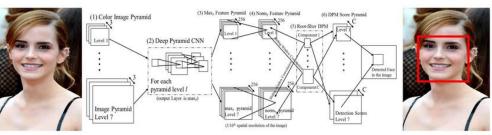
CNN





Applications of CNN

Image recognition Facebook, Apple, MS, etc



Video analysis

Commentary, Self-driving car, etc

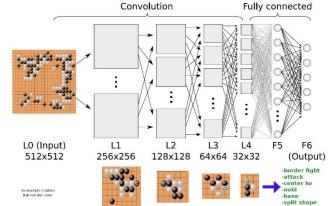
Natural language processing
Translation, automatic captioning, etc

Drug discovery: AtomWise

Games: DeepMind, Go, Atari, etc









Applications of CNN

삼성전자 뛰쳐나와 만든 앱, 찍는 순간 칼로리가 나온다

[출처: 중앙일보] 삼성전자 뛰쳐나와 만든 앱, 찍는 순간 칼로리가 나온다





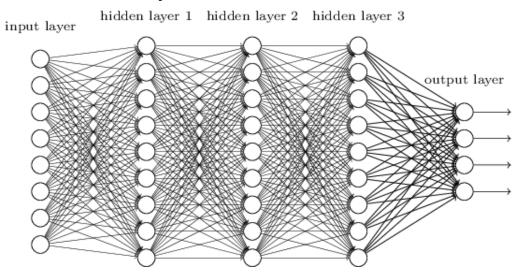


CNN



Drawbacks of DNN

- Previous DNNs (MLPs) use fully-connected layers
- Connect all the neurons between the layers



Drawbacks

Large number of parameters

Easy to be over-fitted
 Large memory consumption (GPU issues)

Does not enforce any structure, e.g., local information



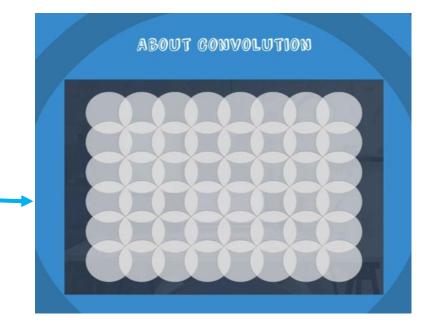
In many applications, local features are important, e.g., images, language, etc.

Key idea of CNN

Weight sharing with a receptive field and convolving it over inputs



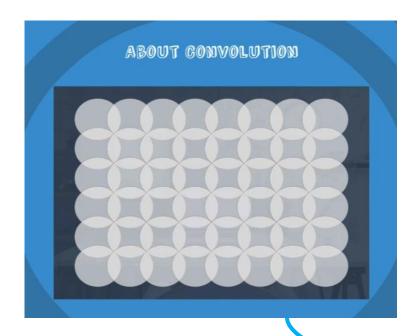
 Each receptive field (filter) shines on specific area of picture and convolve over inputs



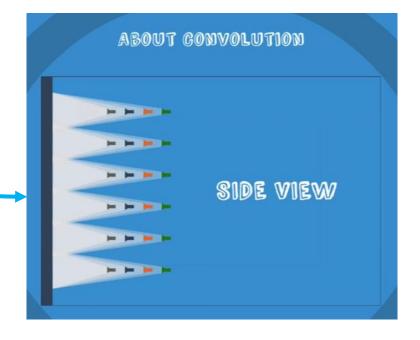


Key idea of CNN

Weight sharing with a receptive field and convolving it over inputs

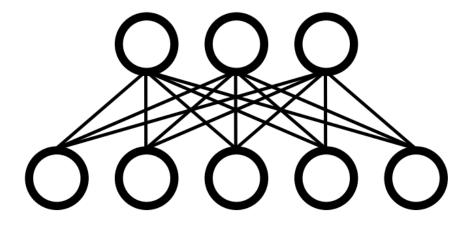


- Each receptive field (filter) shines on specific area of picture and convolve over inputs
- Four different filters shine each area and try to detect patterns.

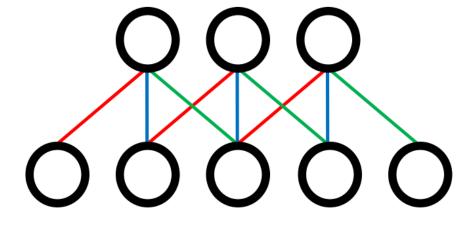


Key idea of CNN

Weight sharing with a receptive field and convolving it over inputs



DNN (fully connected)



CNN (weight sharing and convolving)

- Reduce the number of parameters (less overfitting)
- Learn local features
- Translation invariance



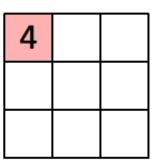
Receptive field & Convolution



Receptive field

Apply same weights over the different spatial regions

1 _{×1}	1,0	1,	0	0
0,0	1,	1,0	1	0
0 _{×1}	0,0	1,	1	1
0	0	1	1	0
0	1	1	0	0



 $h_k^{(l+1)} = \sigma(\sum_{ij \in Fk} w_{ij}^{(l)} h_{ij}^{(l)} + b^{(l)})$

A selection of input maps by the receptive field

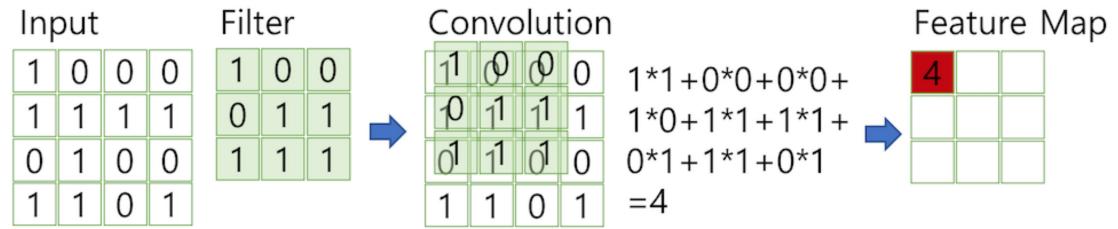
Image

Convolved Feature

Local feature



Receptive field



http://taewan.kim

$$h_k^{(l+1)} = \sigma(\sum_{ij \in Fk} w_{ij}^{(l)} h_{ij}^{(l)} + b^{(l)})$$



Example

Input: 5×5

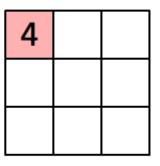
Receptive field (or filter): 3×3

Stride: 1

Pooling: summation

1,	1,0	1,	0	0
0,0	1,	1,0	1	0
0 _{×1}	0,0	1 _{×1}	1	1
0	0	1	1	0
0	1	1	0	0

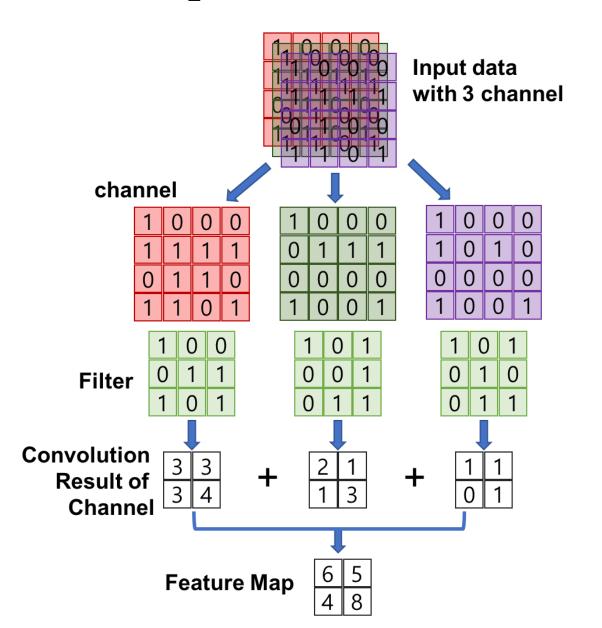
Image



Convolved Feature

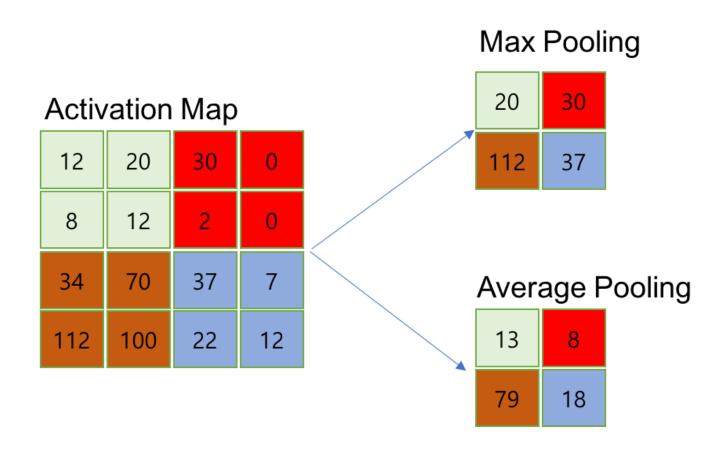


Multiple channels





Pooling

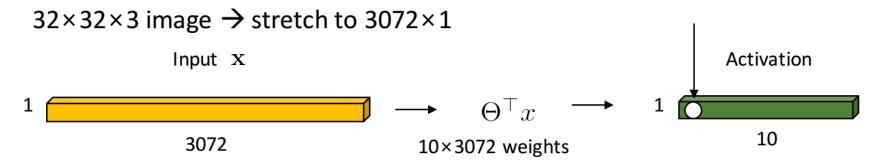


- Make the representations smaller
- Reduce number of parameters and computation
 Regularization effect

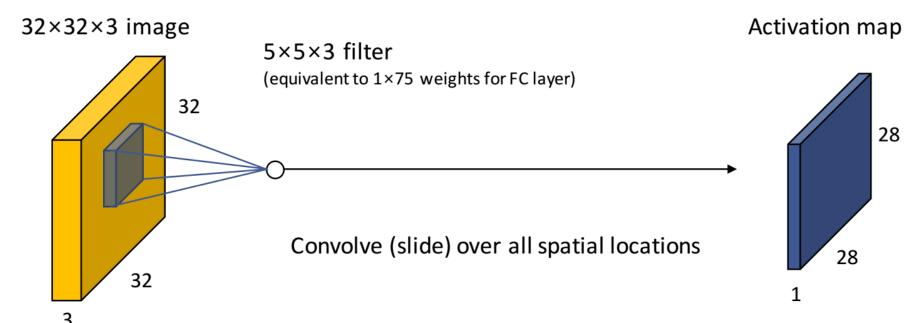


Receptive field

Fully connected layer



Convolution layer

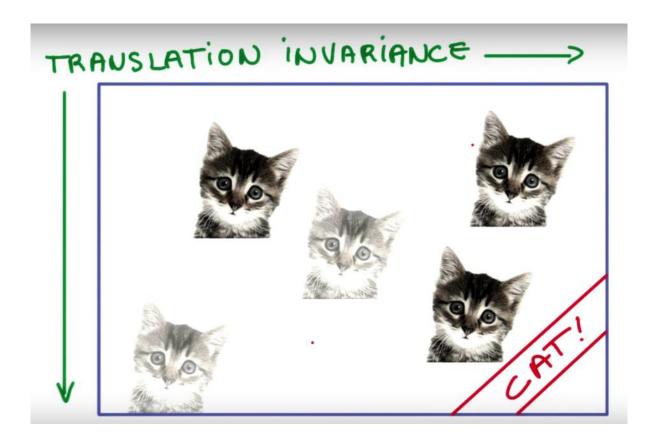




Translation invariance

When input is translated, the corresponding output to recognize the object should not be changed.

CNN can produce the same output even though the input image is shifted due to weight sharing.





Padding



Padding

9

9

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

7×7 input (spatially)
Zero pad 1 pixel border
Assume 3×3 filter
Applied with **stride 3**

 \rightarrow 3×3 output



Formula for output size

9

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Input data hieght: H

Input data width: W

Filter height: FH

Filter width: FW

Stride size: S

Padding size: P

$$OutputHeight = OH = rac{(H + 2P - FH)}{S} + 1$$
 $OutputWeight = OW = rac{(W + 2P - FW)}{S} + 1$



Formula for output size

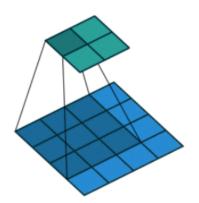
- Output size must be integer
- If a pooling layer coming after conv layer, the size of feature map should be integer times of the pooling size
 - Ex) Pooling size: 3x3
 Then, the output size before the pooling may be 9x9 or 9x6.
- FH, FW, S, P, and Pooling size should be determined in such a way

- Input data hieght: H
- Input data width: W
- Filter height: FH
- Filter width: FW
- Stride size: S
- Padding size: P

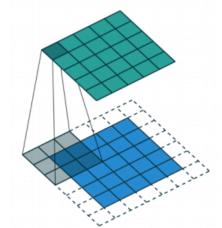
$$OutputHeight = OH = rac{(H + 2P - FH)}{S} + 1$$
 $OutputWeight = OW = rac{(W + 2P - FW)}{S} + 1$



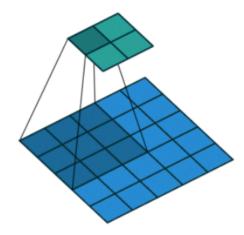
More examples



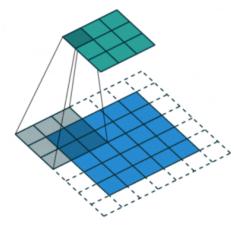
No padding, stride 1



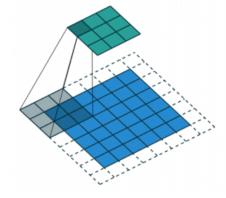
Padding 1, stride 1



No padding, stride 2



Padding 1, stride 2

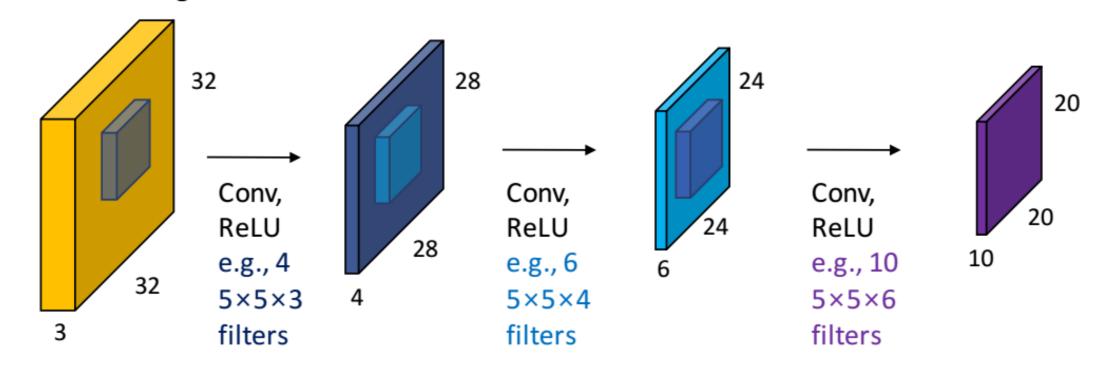


Padding 1, stride 2 (odd)



ConvNet

32×32×3 image





Technical issues

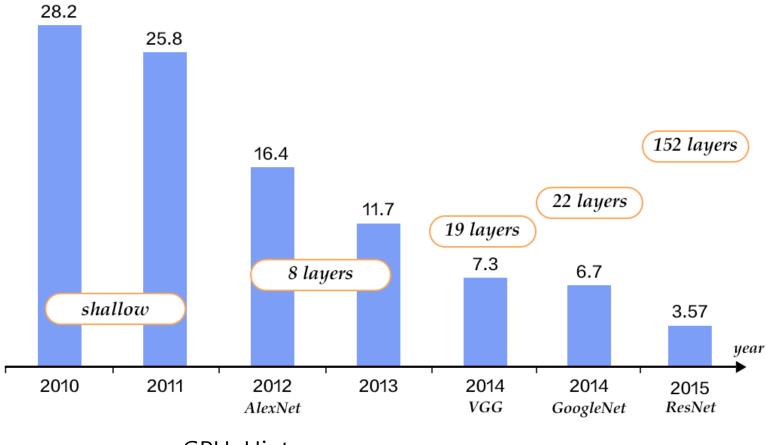
Going wider: inception

Going deeper: skip connection



ILSVRC Winners

ImageNet Large Scale Visual Recognition Challenge

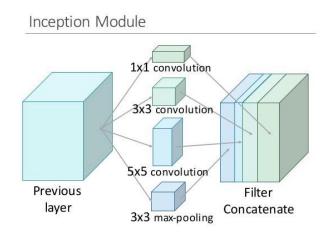




GPU, Hinton group

Inception

- To go wiiiiiider
- To extract different features



Each layer type extracts a different feature.

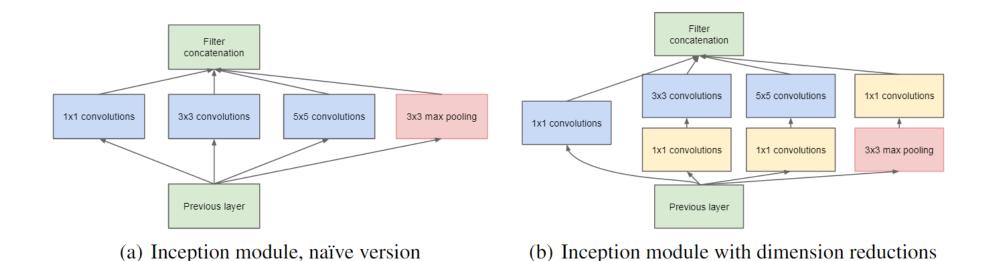
The output of a 5x5 receptive field is different from the output of a 3x3 one.

Computes multiple different transformations over the same input map in parallel.

Then, the network chooses the best combination.



Inception

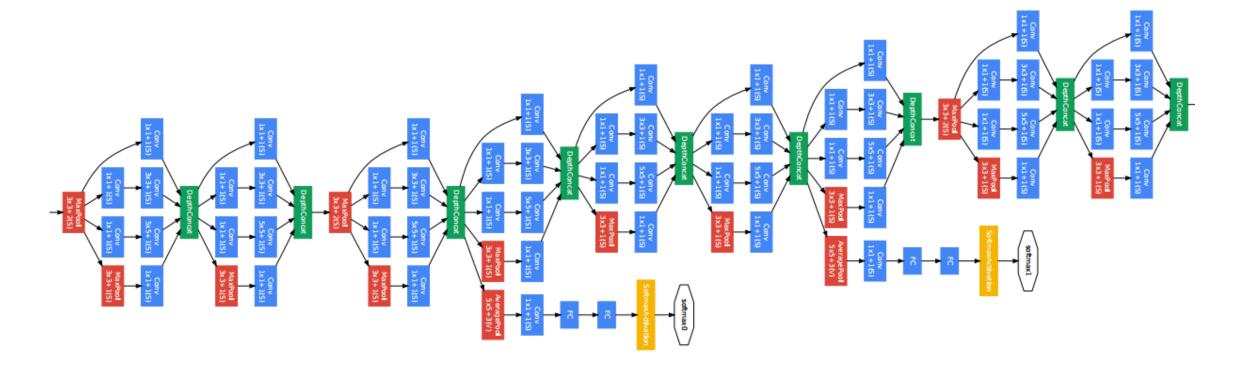


Stack different layer transformations in parallel, resulting in nets that were **simultaneously deep** (many layers) and **wide** (many parallel operations).



Inception

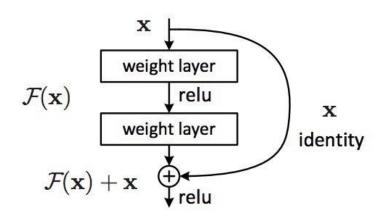
GoogLeNet – Winner of 2014 ImageNet Challenge





Skip connection

- To go deeeeeper
- To avoid the vanishing gradient



Each "block" in ResNet consists of a series of layers and a skip connection adding the input of the block to its output.

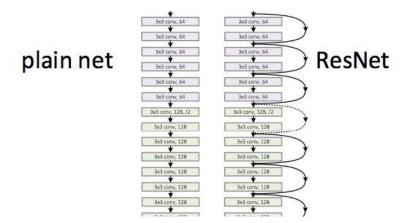
The "add" operation is performed element-wise

If the input and output are of different sizes, zero-padding
or projections (via 1x1 convolutions) can be used to create
matching dimensions.



Skip connection

ResNet (Residual net) – Winner of 2015 ImageNet Challenge



The gradient signal in ResNets could travel back directly to early layers via skip connections

→ Going deeper and deeper to 50-layer, 101-layer, 152-layer, and even 1000+ layer nets.



Examples



NLP example

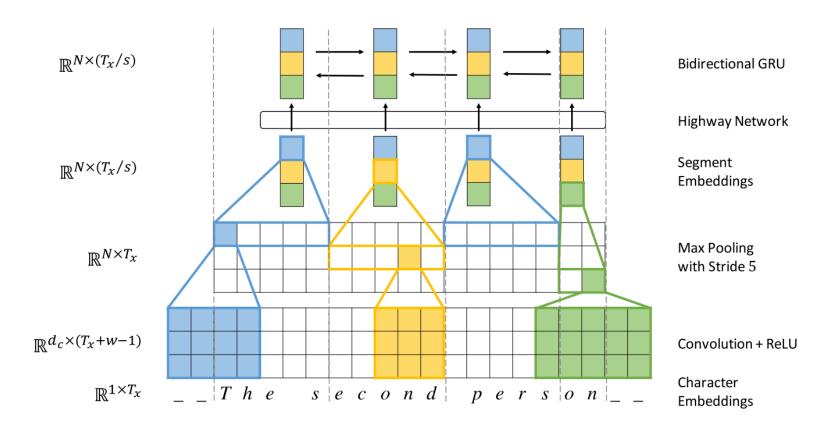


Figure 1: Encoder architecture schematics. Underscore denotes padding. A dotted vertical line delimits each segment.



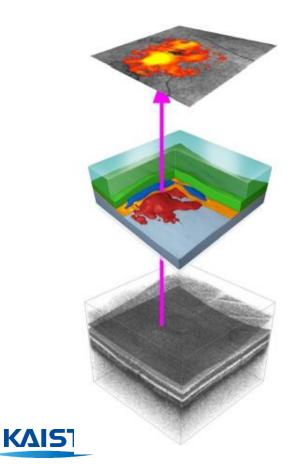
Medical imaging

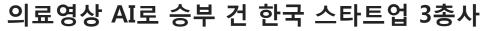
DeepMind





Clinically applicable deep learning for diagnosis and referral in retinal disease



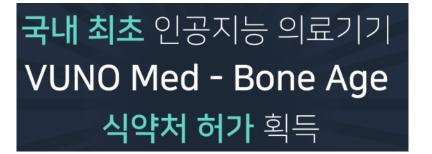










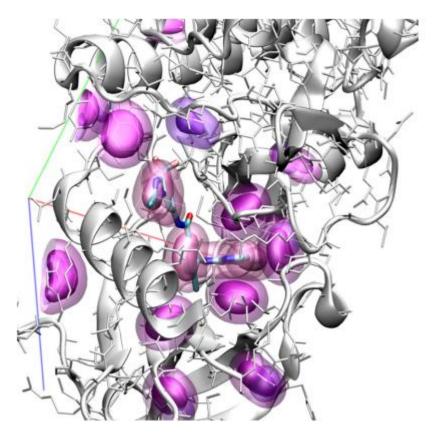


CNN in Chemistry



Protein-Ligand Scoring with Convolutional Neural Networks

Matthew Ragoza,^{†,‡} Joshua Hochuli,^{‡,¶} Elisa Idrobo,[§] Jocelyn Sunseri,[∥] and David Ryan Koes*,[∥]



We describe convolutional neural network (CNN) scoring functions that take as input a comprehensive three-dimensional (3D) representation of a protein—ligand interaction.

Visualization of atom densities used as input to CNN scoring. Aromatic carbon atom densities are shown at two isosurface levels (solid and transparent surfaces) for both the receptor (purple) and ligand (lavender).

our CNN scoring function outperforms the AutoDock Vina scoring function when ranking poses both for pose prediction and virtual screening.

J. Chem. Inf. Model. 2017, 57, 942-957



[†]Department of Neuroscience, [‡]Department of Computer Science, [¶]Department of Biological Sciences, and [∥]Department of Computational and Systems Biology, University of Pittsburgh, Pittsburgh, Pennsylvania 15260, United States

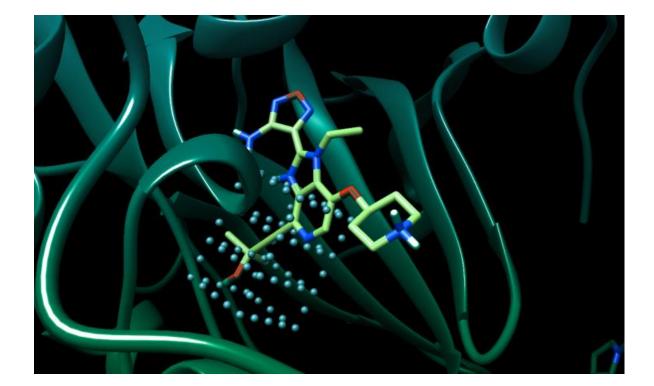
[§]Department of Computer Science, The College of New Jersey, Ewing, New Jersey 08628, United States

AtomNet: A Deep Convolutional Neural Network for Bioactivity Prediction in Structure-based Drug Discovery



Izhar Wallach Atomwise, Inc. izhar@atomwise.com Michael Dzamba Atomwise, Inc. misko@atomwise.com

Abraham Heifets
Atomwise, Inc.
abe@atomwise.com



more than \$51 million in funding





 $L \times L$

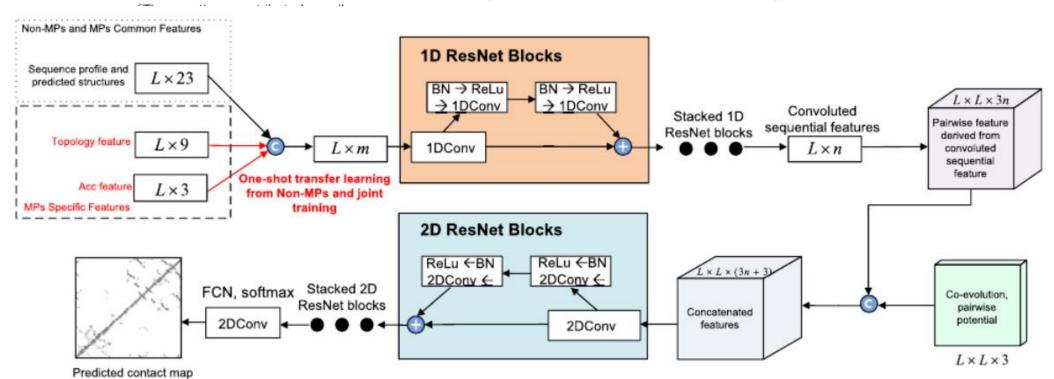


Folding Membrane Proteins by Deep Transfer Learning

Computational elucidation of membrane protein (MP) structures

Sheng Wang, 1,2,4,5 Zhen Li,1,3,5 Yizhou Yu,3 and Jinbo Xu1,6,*

⁴Computational Bioscience Research Center (CBRC), King Abdullah University of Science and Technology (KAUST), Thuwal, Saudi Arabia





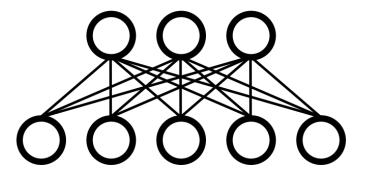
¹Toyota Technological Institute at Chicago, Chicago, IL 60637, USA

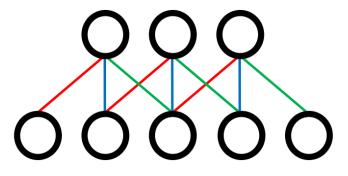
²Department of Human Genetics, University of Chicago, Chicago, IL 60637, USA

³Department of Computer Science, University of Hong Kong, Hong Kong

Summary

Weight sharing with a receptive field and convolving it over inputs



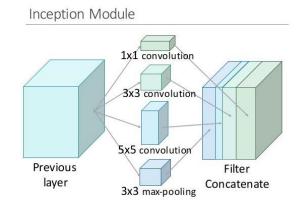


DNN (fully connected)

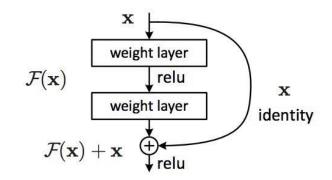
CNN (weight sharing and convolving)

- Reduce the number of parameters (less overfitting)
- Learn local features
- Translation invariance

- To go wiiiiiider
- To extract different features



- To go deeeeeper
- To avoid the vanishing gradient





New terms

- Receptive field
- Convolution
- CNN
- Translation invariance
- Stride
- Pooling
- Feature map
- Padding
- Inception
- Skip connection

